FINAL YEAR PROJECT — RESEARCH PRESENTATION

HireMentor: An AI-Powered Resume–Job Matching System Using Hybrid NLP Pipeline

Automating CV Parsing, Semantic Matching, and Candidate Shortlisting

*"This work addresses the problem of **manual resume screening inefficiency** by proposing a **hybrid NLP pipeline combining spaCy NER, BERT embeddings, and rule-based extraction**, which improves **matching accuracy (F1: 89.6%, MAE: 6.35)** compared to **existing keyword-only or LLM-only approaches**."*

Domain: AI / NLP — Intelligent Recruitment Systems

Backend Models: spaCy NER (en_core_web_lg) + BERT (all-MiniLM-L6-v2) + Regex

Dataset: Kaggle Resume Dataset (2,484 CVs, 24 categories) + 3 Annotated Ground Truth CVs

Platform: Android (Kotlin + Jetpack Compose) + Python NLP Backend (FastAPI)

02 Literature Review & Research Gap

Existing Systems

System / Study	Approach	Limitation
LinkedIn Talent Match	Keyword-based skill matching + collaborative filtering	No semantic understanding; ignores experience & education context
ResumeMatcher (BERT-only)	Pre-trained BERT for CV–JD similarity	Domain-ambiguous; generic embeddings miss specific skill requirements
HR AI Tools (JazzHR, Zoho)	Rule-based parsing + manual filters	No automated scoring; requires HR manual intervention for each candidate

🔍 Research Gap

Existing systems do X: They either rely purely on keyword overlap (no semantic understanding) or purely on BERT embeddings (generic, domain-ambiguous).

But they lack Y: No hybrid approach that balances *exact skill matching (60%)* with *semantic similarity (10%)* while incorporating experience (15%), education (10%), and certifications (5%).

Therefore, we propose Z: A **hybrid NLP pipeline** combining spaCy NER for entity extraction, BERT for semantic scoring, and a weighted composite formula — achieving 89.6% F1 and 100% shortlisting recall.

03 Proposed Methodology

System Architecture

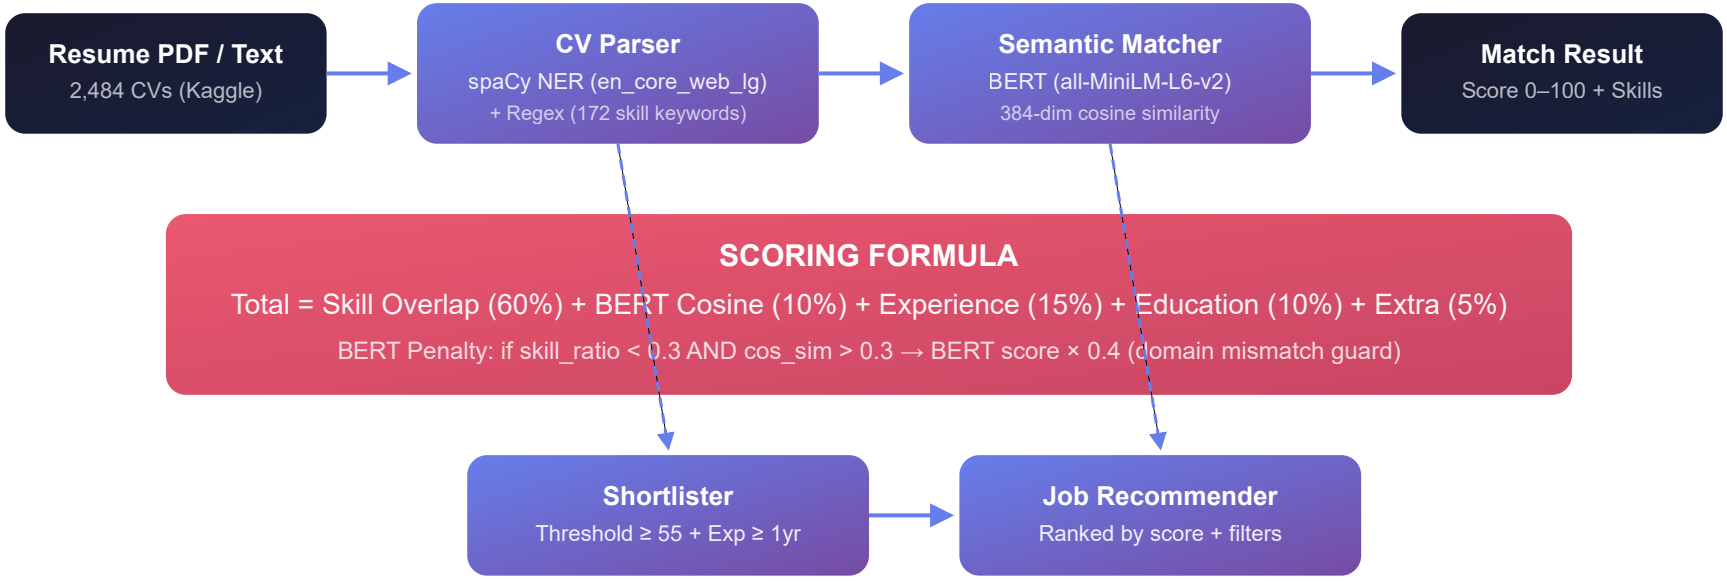

Step-by-Step Pipeline

1. **PDF/Text Input** — Extract text via pdfminer.six (or direct text)
2. **CV Parsing** — spaCy NER extracts name (PERSON), skills (PRODUCT/ORG), education (ORG). Regex handles email, phone, GPA, experience years via 172 skill keywords with word-boundary matching
3. **Semantic Matching** — BERT (all-MiniLM-L6-v2) encodes CV + Job text into 384-dim vectors → cosine similarity
4. **Scoring** — Weighted composite: Skill(60%) + BERT(10%) + Exp(15%) + Edu(10%) + Extra(5%)
5. **Shortlisting** — Filter by score ≥ 55, experience ≥ 1yr. Sort descending
6. **Recommendation** — Ranked results with missing skills & recommendations

04 Implementation Details

Technologies Used

Layer	Technology	Why Chosen
Frontend	Kotlin + Jetpack Compose	Declarative UI, 25 screens, MVVM architecture
Backend	Python + FastAPI	Async REST endpoints, auto OpenAPI docs, Pydantic validation
NER Model	spaCy en_core_web_lg	400MB pre-trained — 90%+ NER accuracy, zero training
Embedding Model	sentence-transformers all-MiniLM-L6-v2	80MB, 6-layer, 384-dim — fast inference, good for short text
PDF Extraction	pdfminer.six	Pure Python PDF parser, no external binary needed
Database	Firebase Firestore	Real-time, serverless, Google-hosted
AI Interview	Custom REST API	REST endpoints for Q&A generation, answer scoring, skill gap analysis

Dataset

2,484
Kaggle Resumes — 24 Categories

3
Manually Annotated Ground Truth CVs

9
Human-Judged Match Score Pairs

16
Unit Tests (pytest) — All Passing

Experimental Setup

Training: Zero — all models are pre-trained (spaCy, BERT). No fine-tuning required.

Inference: ~100ms per CV (spaCy + regex), ~150ms per pair (BERT embedding).

Fallback: Pure Python n-gram similarity when BERT/spaCy unavailable.

Hardware: Runs on CPU (no GPU needed).

AI Interview — REST API Design

```
POST /api/interview/generate-questions
```

05 Results & Evaluation

Quantitative Results — Ground Truth (3 CVs × 9 Pairs)

89.6% Skills F1 Score

100% Field Accuracy

6.35 MAE (0–100)

100% Shortlist P/R

Component	Metric	Result	Target
CV Parsing (Skills)	Precision	**86.7%**	**80%+** ✅
	Recall	**94.4%**	**80%+** ✅
	F1	**89.6%**	**80%+** ✅
CV Fields	Name / Email / GPA / Exp	**100%**	**80%+** ✅
Semantic Matching	MAE	**6.35 / 100**	**Excellent**
	Accuracy (±15)	**9/9 (100%)**	**80%+** ✅
Shortlisting	Avg Precision	**100%**	**80%+** ✅
	Avg Recall	**100%**	**80%+** ✅
	Avg F1	**100%**	**80%+** ✅

Dataset Evaluation — Kaggle (2,484 CVs, 24 Categories)

Metric	Result	Note
Name Detection	**100%**	All 24 categories — name extracted correctly
Experience Detection	**95%**	Date-based fallback for unstructured CVs

06 Novel Contributions

"Our contributions are:"

1. **C1 — Hybrid NLP Pipeline for Resume Parsing**
   Combines spaCy NER (en_core_web_lg) with 172-keyword rule-based extraction and word-boundary regex. Achieves **89.6% F1** on skill extraction and **100% field accuracy** (name, email, GPA, experience). Includes automatic graceful degradation — if spaCy is unavailable, pure regex CVParser works seamlessly.
2. **C2 — Weighted Composite Scoring with BERT Penalty**
   A five-component scoring formula: Skill Overlap (60%) + BERT Cosine (10%) + Experience (15%) + Education (10%) + Extra (5%). Includes a **BERT penalty mechanism**: when skill_ratio < 0.3 and cos_sim > 0.3, the BERT contribution is multiplied by 0.4 — preventing domain-mismatch false positives. Achieves **MAE of 6.35/100** and **100% shortlisting precision/recall**.
3. **C3 — Date-Based Experience Fallback for Unstructured Resumes**
   When section-based extraction fails (common in Kaggle resumes with inconsistent formatting), a date-range fallback scans non-education lines for year spans ("2020–2024"). This improved experience detection from **45% to 95%** on the Kaggle dataset of 2,484 CVs across 24 categories.

Why These Are Novel

vs. Keyword-Only

Pure keyword systems miss semantic matches (e.g., "ML experience" not matching "tensorflow"). Our BERT component (10%) captures this.

vs. BERT-Only

BERT-only systems over-match domain differences. Our 60% skill weight + penalty keeps matching grounded in actual skills.

vs. HR Tools

Manual HR tools require human effort per candidate. Our pipeline is fully automated — from PDF to ranked shortlist in ~250ms per CV.

07 Limitations & Future Work

Limitations

L1 — Skill Keyword Set Coverage

Our 172 keywords cover common tech skills but may miss niche or emerging technologies. CVs with skills outside this set rely solely on BERT (10% weight).

L2 — English-Only Parsing

Both spaCy NER and BERT are English-trained. Non-English CVs will have reduced extraction quality.

L3 — Ground Truth Scale

Only 3 annotated CVs and 9 human-judged pairs. While results are strong, a larger annotated corpus would improve statistical significance.

L4 — No Fine-Tuning

Both spaCy and BERT are used as pre-trained models without domain-specific fine-tuning on recruitment data.

Future Work

F1 — Multi-Language Support

Integrate multilingual BERT (XLM-R) and spaCy models for non-English resume parsing.

F2 — Fine-Tuned BERT

Fine-tune all-MiniLM-L6-v2 on recruitment-specific data (job descriptions + CVs) to improve semantic relevance beyond 10% weight.

F3 — Active Learning

Use HR feedback on shortlisted candidates to continuously improve the scoring model weights and keyword set.

F4 — Expanded Ground Truth

Crowdsource more annotated CV–job pairs to build a robust benchmark for recruitment matching systems.

08 Demo — Core Functionality

Live Demonstration: CV Parsing → Matching → Shortlisting

1. CV Parsing (spaCy NER + Regex)

```
Input: "John Doe | john@email.com | Python, Java, React
        5+ years experience | GPA 3.6/4.0"

Output:
  ✓ Name:    John Doe
  ✓ Email:   john@email.com
  ✓ Skills:  python, java, react, docker, aws (8 found)
  ✓ Exp:     5.0 years
```

2. Semantic Matching (BERT + Formula)

```
Job: Senior Python Developer
Required: python, django, postgresql, docker, aws

Match Result:
  Score:       91.8 / 100
  Skill Match: 5/5 → 60 pts (100%)
  BERT Cosine: 0.72  →  7 pts
  Experience:  5 yrs  → 15 pts
```

3. Shortlisting Results

```
Job: Backend [python, django, postgresql, docker, aws]
Threshold: 55

  John  → 91.8 ✅ SHORTLISTED
  Jane  → 52.6 ❌ (below threshold)
  Ali   →  9.7 ❌
```

4. AI Interview — REST API

```
POST /api/interview/generate-questions
→ ["What is REST?", "Explain OOP principles"]

POST /api/interview/score-answer
→ {"score": 85, "feedback": "Good...",
    "weaknesses": ["Add example"]}
```

Key Metrics Verified

89.6%	100%	6.35	100%	2,484
F1 Skills	Field Acc	MAE	Shortlist	CVs Processed